

Estimation of Dynamic Discrete Demand Model with Bayesian Markov Chain Monte Carlo

Bayesian Estimation of Dynamic Discrete Choice Models

Author(s): Susumu Imai, Neelam Jain and Andrew Ching

Source: *Econometrica*, Vol. 77, No. 6 (Nov., 2009), pp. 1865-1899


Published by: The Econometric Society

Raymond Chiasson

RSM3055

ICJ Method Assignment

8-Nov-2022



Summary

I solve the Loyalty Program model using the IJC method in Julia. These slides briefly document the model, method, code, and estimation results.

This project features

- Bayesian Markov Chain Monte Carlo (MCMC), which is an efficient method to estimate structural models with many parameters such as those with heterogeneous individuals.
- The “IJC method” described in Imai, Jain and Ching 2009 which is a computationally efficient way to approximate the value function in dynamic models which are being estimated with Bayesian MCMC.
- The “Loyalty Program” which is a basic discrete dynamic choice model where consumers decide between products in the current period, and can earn “Rewards” after having chose a product a certain number of times (for example, a “buy 5 get one free”).
- Clean and optimized Julia code using best practices of the language, to the best of my knowledge.

Model Parameters and Environmental Variables

```
# -----  
# 0. Parameters and setup.  
# -----  
# Model parameters (to be estimated).  
 $\alpha$  = [-0.0, -0.0, -0.0]      # store brand intercepts  
 $\gamma$  = -1.0                  # price coefficient  
G = [1.0, 3.0, 8.0]          # value of gifts  
 $\sigma_1$  = [0.0, 0.0, 0.0]    # homogeneous consumers  
 $\beta$  = 0.650                 # discount rate  
  
# Environment variables.  
n_stores = 3                 # number of stores  
n_choices = n_stores + 1    # number of choices for consumer (stores + outside option)  
 $\bar{s}$  = [4, 4, 6]            # gift threshold  
price_mean = [1.0, 0.75, 1.5] # mean of observed prices  
price_stdev = [0.25 0.00 0.00;  
               0.00 0.25 0.00;  
               0.00 0.00 0.25] # standard deviation (covariance matrix) of observed prices  
n_price_draws = 100          # number of draws for price integration  
  
# Simulated data parameters.  
n_individuals = 1000  
n_periods = 100
```

Parameter Estimates – Bayesian MCMC with IJC

Bayesian MCMC with IJC Results

8x5 DataFrame

Row	variable	true value	estimated mean	95% Credible Interval Low	95% Credible Interval High
	Any	Any	Any	Any	Any
1	α_1	-0.0	-0.0106112	-0.0542716	0.0367328
2	α_2	-0.0	-0.0207863	-0.0559575	0.0169452
3	α_3	-0.0	0.00135663	-0.0615654	0.0694388
4	γ	-1.0	-0.997628	-1.04115	-0.957486
5	G_1	0.65	0.654841	0.641717	0.667859
6	G_2	1.0	1.01293	0.971786	1.05542
7	G_3	3.0	3.1305	3.07501	3.18438
8	β	8.0	7.89058	7.70638	8.04987

completed 30000 draws, acceptance rate 0.24567408875385913

Parameter Estimates – Maximum Likelihood

Numerical Maximum Likelihood Results					
8x5 DataFrame					
Row	values	true value	estimated mean	95% Confidence Interval Low	95% Condifence Interval High
	Any	Any	Any	Any	Any
1	α_1	-0.0	-0.0061304	-0.0491608	0.0369
2	α_2	-0.0	0.000857478	-0.0365915	0.0383065
3	α_3	-0.0	-0.000520678	-0.0618666	0.0608253
4	γ	-1.0	-0.998529	-1.03678	-0.960279
5	G_1	1.0	0.995541	0.930543	1.06054
6	G_2	3.0	3.02741	2.93073	3.12408
7	G_3	8.0	7.9929	7.83024	8.15557
8	β	0.65	0.650514	0.639755	0.661272

Initial Parameter Guesses

Variable	Parameter Guess			
	True Value	Informed	Directional	Zeros
α_1	0	0	1	0
α_2	0	0	1	0
α_3	0	0	1	0
γ	-1	-0.5	-1	0
β	0.65	0.8	0.5	0
G_1	1	2	1	0
G_2	3	5	1	0
G_3	8	5	1	0



I experimented with different starting parameter values.

- Starting from the true value, to confirm that the distribution converged around these values.
- Informed and directional guesses which were logical choices based on economic theory and preliminary runs.
- A naïve guess of 0 for each parameter.



The distribution converged well before the burn-in period for all parameter guesses.

The following graphs and results all use the Zeros parameter guess.



Accepted Parameter Draws from Bayesian Posterior




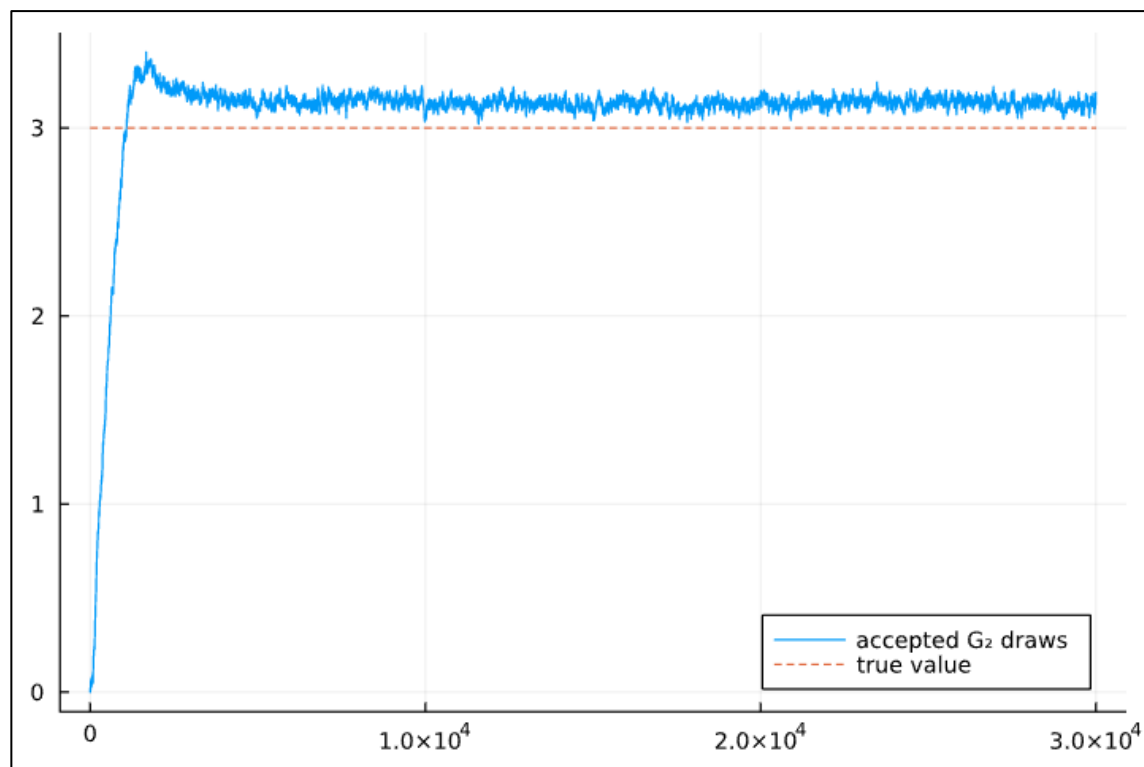
Accepted Parameter Draws from Bayesian Posterior



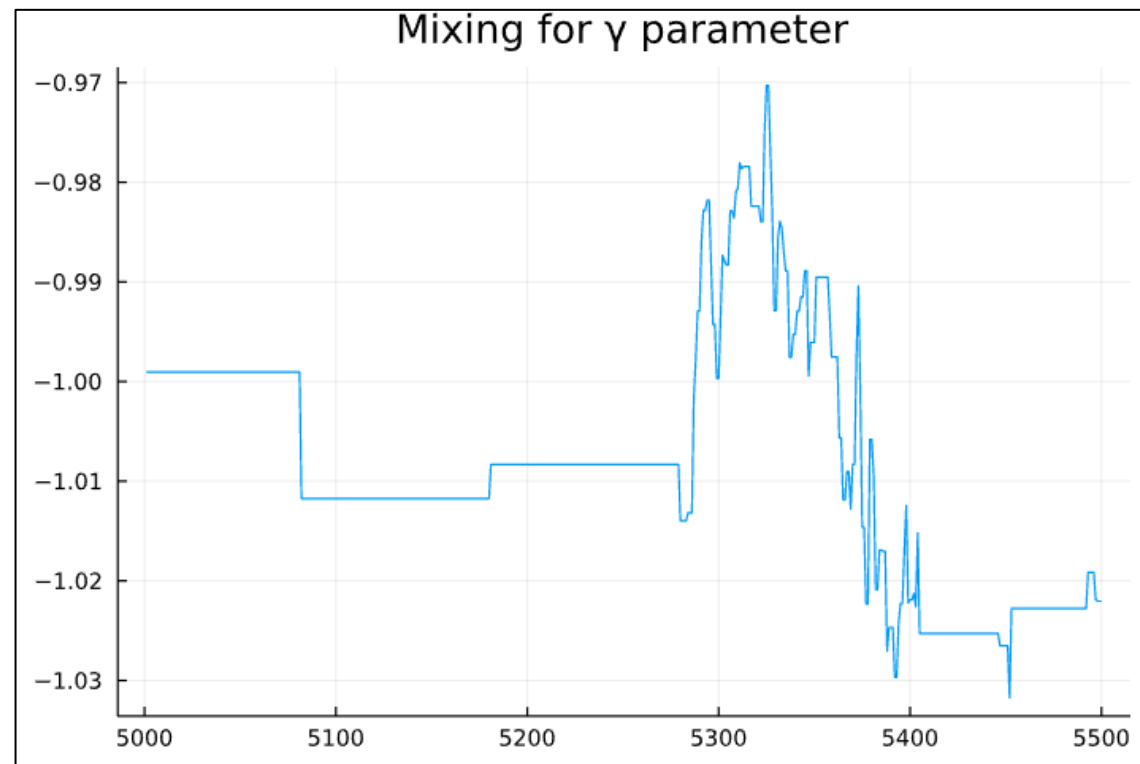

Accepted Parameter Draws from Bayesian Posterior




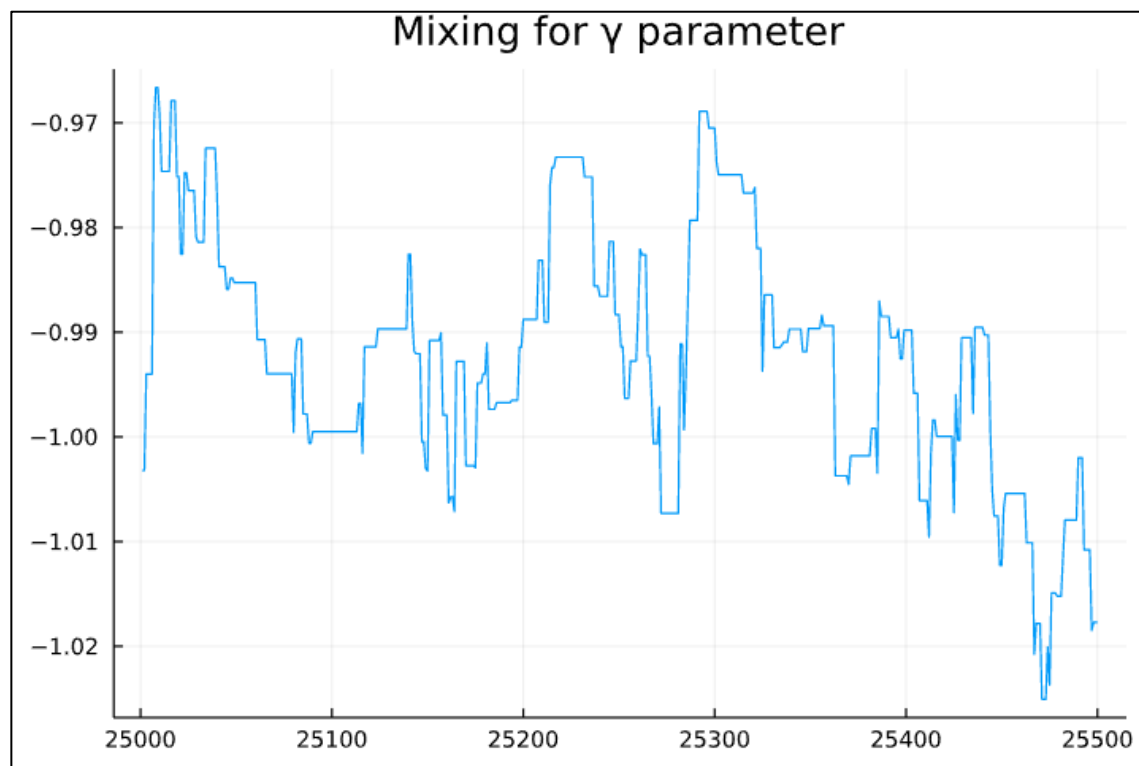
Accepted Parameter Draws from Bayesian Posterior




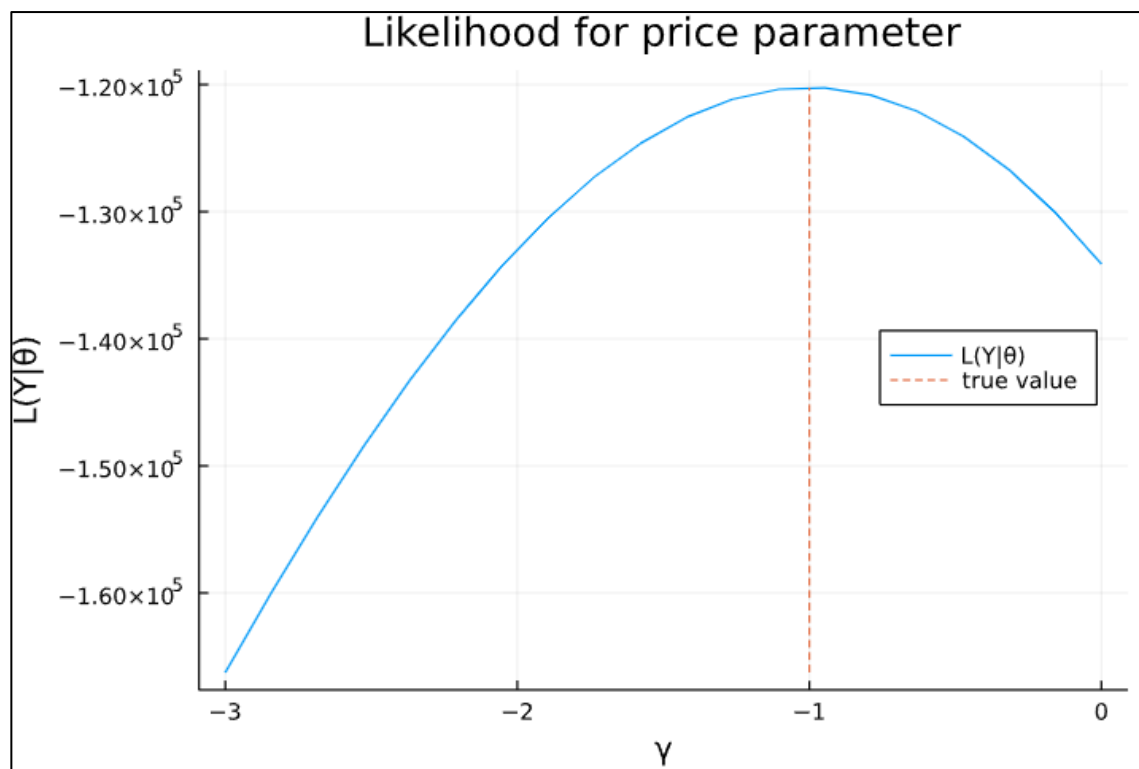
Parameter Draw Mixing During Burn-In



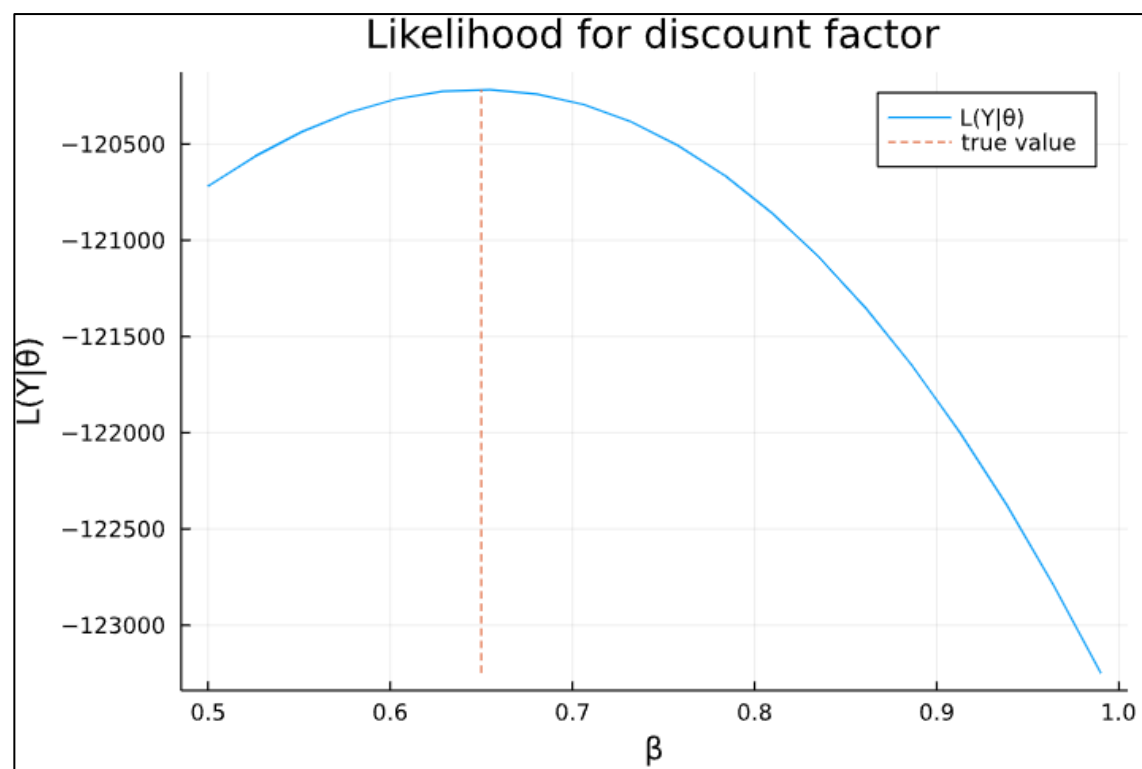
Parameter Draw Mixing After Burn-In



Sample Log Likelihood Function



Sample Log Likelihood Function



Model Parameters and Environmental Variables

```
# -----  
# 0. Parameters and setup.  
# -----  
# Model parameters (to be esimated).  
 $\alpha$  = [-0.0, -0.0, -0.0]          # store brand intercepts  
 $\gamma$  = -1.0                      # price coefficient  
G = [1.0, 3.0, 8.0]              # value of gifts  
 $\sigma_g$  = diagm([0.5, 1.0, 2.0])  # covariance of gift values (heterogeneous consumers)  
 $\beta$  = 0.650                     # discount rate  
  
# Environment variables.  
n_stores = 3                     # number of stores  
n_choices = n_stores + 1         # number of choices for consumer (stores + outside option)  
 $\bar{s}$  = [4, 4, 6]                 # gift threshold  
price_mean = [1.0, 0.75, 1.5]    # mean of observed prices  
price_stdev = diagm([0.25, 0.25, 0.25]) # standard deviation (covariance matrix) of observed prices  
n_price_draws = 100              # number of draws for price integration  
  
# Simulated data parameters.  
n_individuals = 100  
n_periods = 1000  
  
# Take draws for individual-level parameters  
 $G_i$  = rand(MvNormal(G,  $\sigma_g$ ), n_individuals)
```

Parameter Estimates

Row	variable	true value	estimated mean	95% Credible Interval Low	95% Credible Interval High
	Any	Any	Any	Any	Any
1	α_1	-0.0	-0.036922	-0.0876735	0.0119379
2	α_2	-0.0	-0.0379751	-0.0778946	0.00114623
3	α_3	-0.0	-0.0632642	-0.13365	0.00545934
4	γ	-1.0	-0.970838	-1.01471	-0.925212
5	β	0.65	0.658644	0.648444	0.668853
6	\bar{G}_1	1.0	1.12314	0.967387	1.28021
7	\bar{G}_2	3.0	3.04384	2.81318	3.27425
8	\bar{G}_3	8.0	7.93703	7.62853	8.25045
9	σ_{g1}	0.5	0.537511	0.394034	0.7284
10	σ_{g2}	1.0	1.22921	0.908963	1.64703
11	σ_{g3}	2.0	1.83767	1.29399	2.55138
12	G_2-10	4.07	3.52406	2.97044	4.10188
13	G_2-50	3.91	3.99739	3.40701	4.60856
14	G_2-100	3.62	3.67842	3.11153	4.24063

completed 30000 draws, homogeneous parameter acceptance rate 0.27
heterogeneous: [0.32 0.31 0.31 0.33 0.34 0.34 0.31 0.35 0.33 0.33]

Initial Parameter Guesses

Variable	Parameter Guess			
	True Value	Informed	Directional	Zeros
α_1	0	0	0	0
α_2	0	0	0	0
α_3	0	0	0	0
γ	-1	-1	-1	0
β	0.65	0	0	0
G_1	1	4	1	0
G_2	3	4	1	0
G_3	8	4	1	0
α_1	0.5	1	1	1
α_2	1	1	1	1
α_3	2	1	1	1



I experimented with different starting parameter values.

- Starting from the true value, to confirm that the distribution converged around these values.
- Informed and directional guesses which were logical choices based on economic theory and preliminary runs.
- A naïve guess of 0 for each parameter.



The distribution converged well before the burn-in period for all parameter guesses.

The following graphs and results all use the Zeros parameter guess.



Accepted Parameter Draws from Bayesian Posterior





Accepted Parameter Draws from Bayesian Posterior





Accepted Parameter Draws from Bayesian Posterior



Accepted Parameter Draws from Bayesian Posterior





Accepted Parameter Draws from Bayesian Posterior




Sample Log Likelihood Function



Heterogenous Case



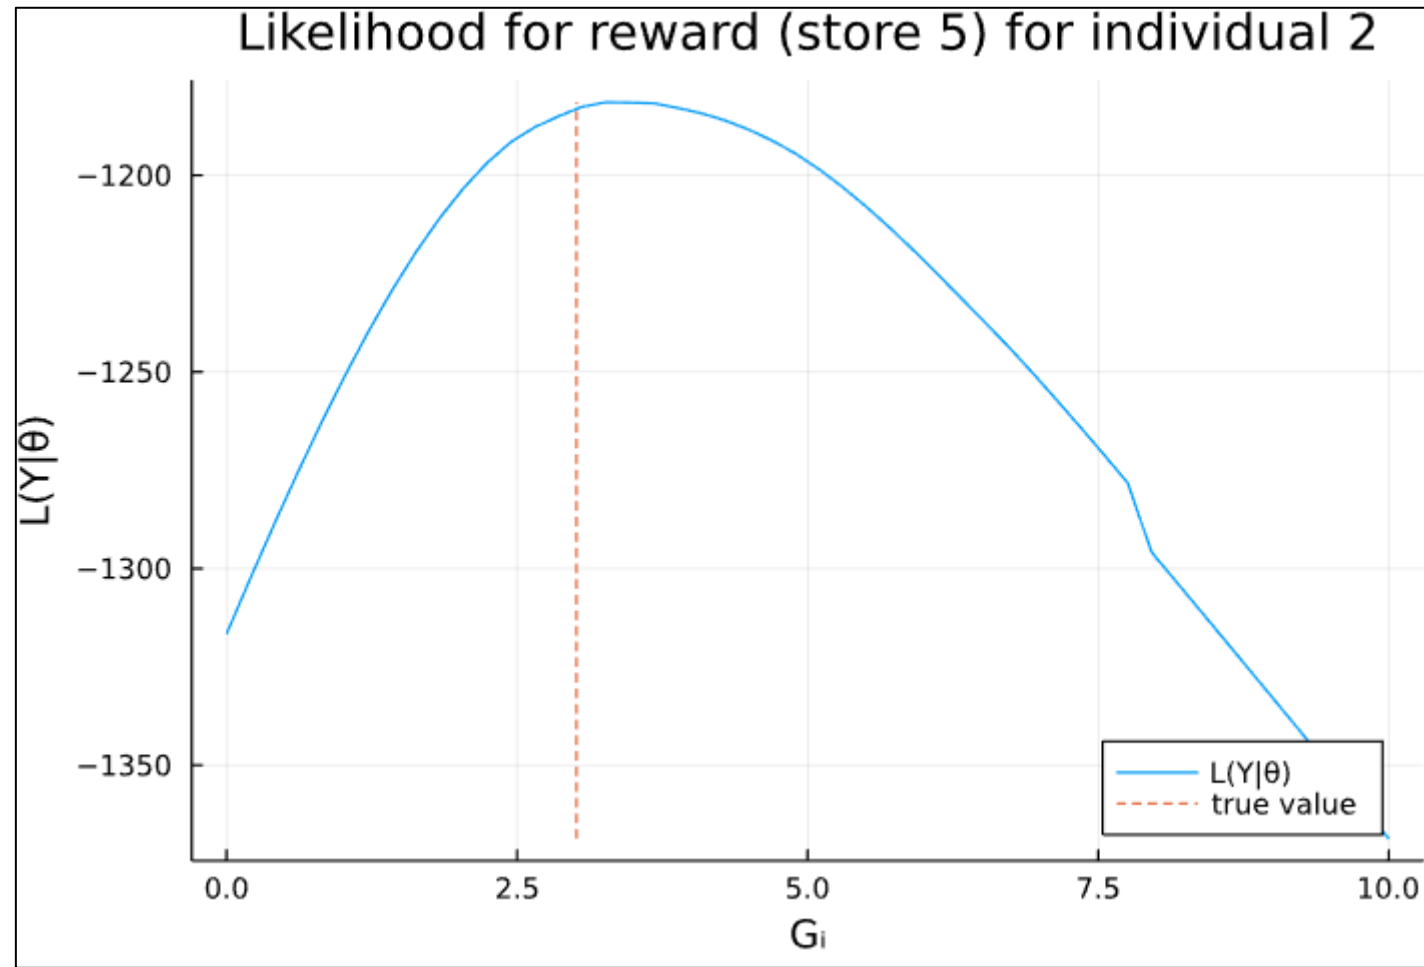
Sample Log Likelihood Function



Sample Log Likelihood Function



Sample Log Likelihood Function




Issue with the Pseudo Value Function Approximation

The pseudo value function approximation to the true value function leads the likelihood function to be non-smooth and biased towards the population mean.

The bias towards the population mean causes incorrect estimates for the variance parameters.

The smoothness can be addressed by adjusting the kernel bandwidth parameter.

The bias can be addressed by using more periods (such as 1000 instead of 100) and by storing individual pseudo value function approximations for each individual (instead of the combined value function approach that stores a random individual in each period).



Outline of Code Files

File	Key sections	Notes
main.jl	<ol style="list-style-type: none"> Parameters and setup Simulate data Estimate model using IJC algorithm Display results Estimate model using Maximum Likelihood Benchmarking, Profiling, Visualizing 	<p>Sections 0-3 perform the key steps of creating the dataset, drawing from the posterior using the IJC algorithm, and displaying the results.</p> <p>Section 4 is only present in the homogeneous case and is shown only for comparison.</p> <p>Section 5 contains additional information used for debugging and optimization.</p>
custom_functions.jl	<ol style="list-style-type: none"> Precalculate_State_Correspondence Solve_Individual_Problem Solve_Value_Function Get_Proposal Parameter_Distance_Kernel Log_Prior_Individual_Params Log_Sample_Likelihood Get_Pseudo_Value_Function 	<p>Function 1 precalculates which states are accessible from each starting state and the corresponding choices and rewards.</p> <p>Function 2 is used to perform one iteration of value function iteration.</p> <p>Functions 4-8 are used when taking parameter draws.</p>

File	Key sections	Notes
data_structs.jl	<p>Model_Parameters_Struct</p> <p>Environment_Struct</p> <p>Dataset_Struct</p> <p>Value_Function_Parameters_Struct</p>	<p>Immutable data structs to store scalars and arrays used by various functions and steps in the main program. Model parameters refers to the key unobservables: α, γ, β etc.</p> <p>Environment struct contains observable model parameters such as the number of stores and the number of visits to earn a reward.</p> <p>Dataset struct contains the simulated dataset; store choices, starting states, and prices.</p> <p>The value function struct contains information required to calculate the pseudo value function; past value functions and corresponding parameter draws.</p>

Optimization

The majority of compute time is spent on unavoidable log and exponential calculations.

- The functions used to update the value function (Solve_Individual_Problem) and calculate the sample log likelihood (Log_Sample_Likelihood) were the computational bottlenecks in the Homogeneous Case.
- Both functions were optimized such that they were bottlenecked by calls to the log and exponential functions.
- The Heterogeneous Case was bottlenecked by estimating the value function (Get_Pseudo_Value_Function). This function was further optimized in this set of code to also be limited by basic math operations.
- The fact that estimating the value function becomes the bottleneck illustrates the utility of the IJC Method, especially in the case of heterogeneous individuals.

Parallelization is employed in the Heterogenous Case to greatly improve the speed of the computation bottleneck (calculating the pseudo value function).

- Calculating the log likelihood can be parallelized over individuals for a given parameter draw.
- Parallelization did not increase the speed of the Homogeneous Case. Calculating the likelihood for each individual did not require estimating the pseudo value function which resulted in the overhead of multithreading causing a net loss in performance.

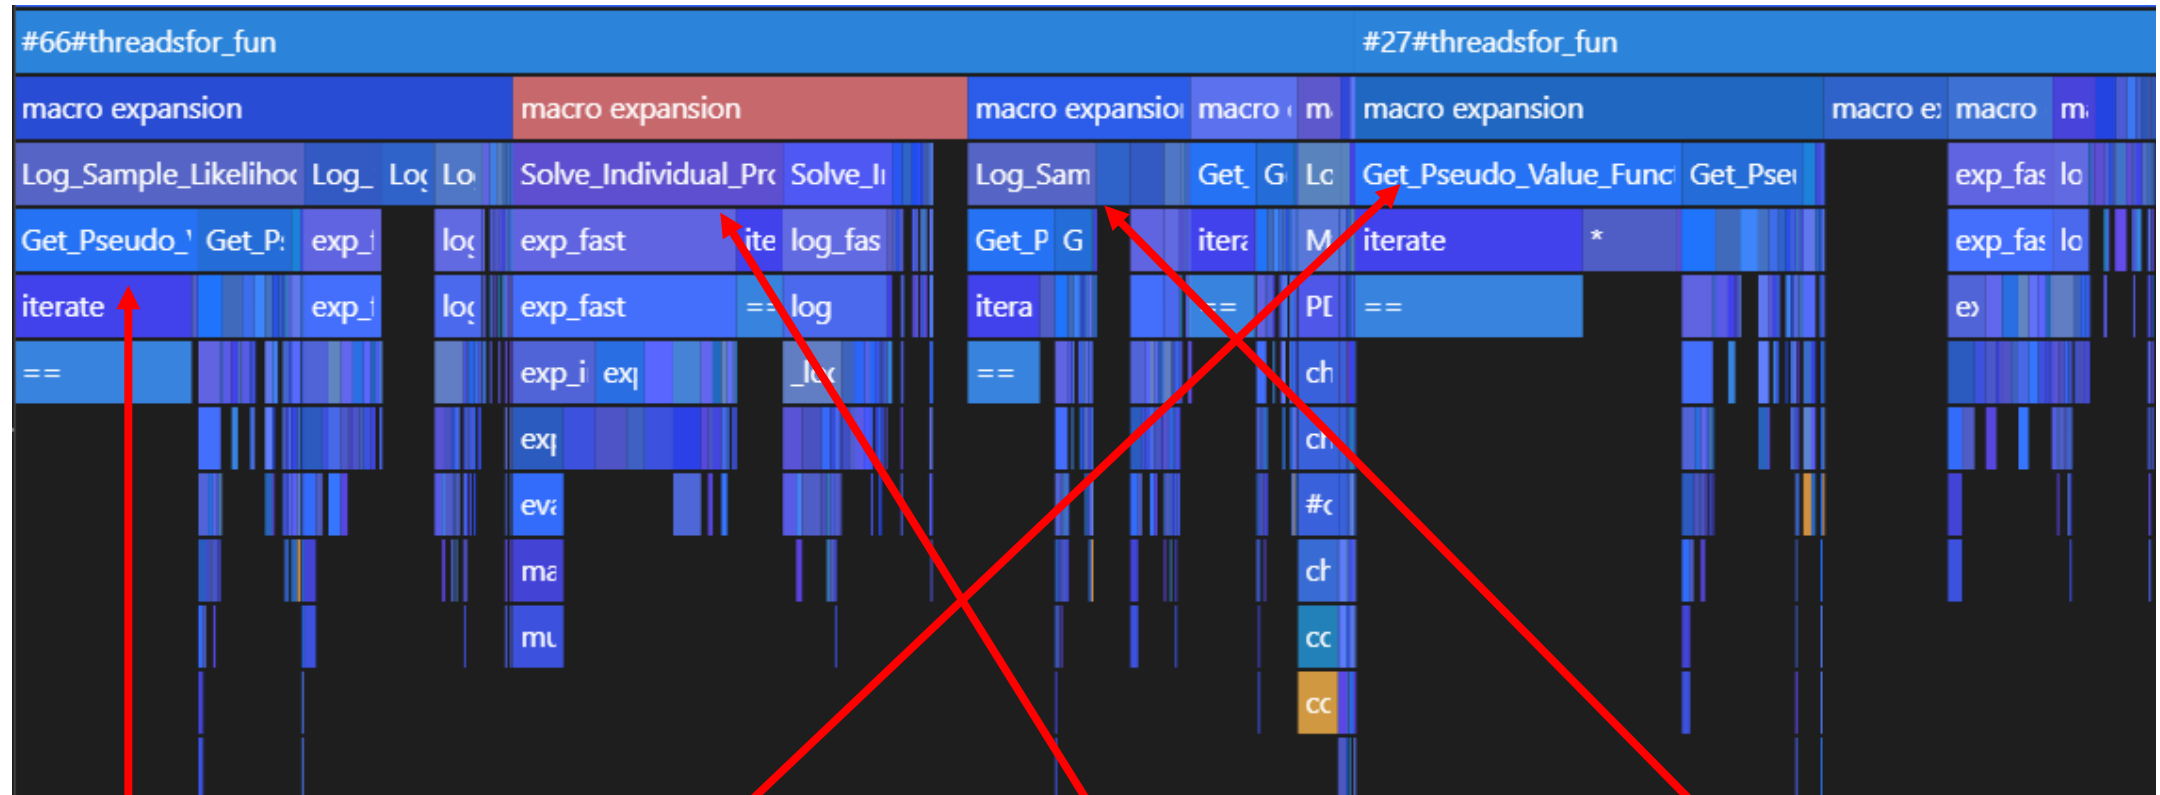
Optimization

Model	Dataset	Method	Compute Time
Homogeneous individuals	1000 individuals 100 periods	IJC method - Bayesian MCMC	15 minutes
		Bayesian MCMC without IJC (solving value function instead)	54 minutes
		Maximum Likelihood with numerical gradients	2.5 minutes
Heterogeneous individuals	1000 individuals 100 periods	IJC method - Bayesian MCMC	290 minutes
		Bayesian MCMC without IJC (solving value function instead)	792 minutes
	100 individuals 1000 periods	IJC method - Bayesian MCMC	23 minutes
		Bayesian MCMC without IJC (solving value function instead)	786 minutes

Optimization

L2										L2									
Solve_Value_Function										L	L	L							
Solve_Individual_Problem					Solve_Individual_Problem					Sol	Sc		exp_fast	log_fast					
exp_fast					itera	log_fast							exp_fast	log					
exp_fast					==	log							exp_	ex	_log				
exp_im	exp_in	ex	e)			_log	_lo	_l					ex						
expm	rei	+	&	>		lo	lo	-					ev						
eval			<			*							m:						
mac													mi						
mul																			

Profiling (IJC method, individual value functions)



Get_Pseudo_Value_Function: 36%

Solve_Individual_Problem: 15%

Log_Likelihood: 10%